

Dealing with Missing Values

```
In [1]: # import libraries
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
```

```
In [2]: # Load the data
data = sns.load_dataset('titanic')
data.head()
```

```
Out[2]:
```


	survived	pclass	sex	age	sibsp	parch	fare	embarked	class	who	adult_male	deck	embark_town	alive	alone
0	0	3	male	22.0	1	0	7.2500	S	Third	man	True	NaN	Southampton	no	False
1	1	1	female	38.0	1	0	71.2833	C	First	woman	False	C	Cherbourg	yes	False
2	1	3	female	26.0	0	0	7.9250	S	Third	woman	False	NaN	Southampton	yes	True
3	1	1	female	35.0	1	0	53.1000	S	First	woman	False	C	Southampton	yes	False
4	0	3	male	35.0	0	0	8.0500	S	Third	man	True	NaN	Southampton	no	True

```
In [3]: # Import libraries
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns

# Load titanic dataset
data = sns.load_dataset('titanic')

# Visualize the data
plt.figure(figsize=(8, 5))
```

```
sns.heatmap(data.isnull(), cbar=False)  
plt.show()
```



```
In [4]: data.info()
```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 891 entries, 0 to 890
Data columns (total 15 columns):
 #   Column        Non-Null Count  Dtype  
---  -
 0   survived      891 non-null    int64  
 1   pclass        891 non-null    int64  
 2   sex           891 non-null    object  
 3   age           714 non-null    float64 
 4   sibsp         891 non-null    int64  
 5   parch         891 non-null    int64  
 6   fare          891 non-null    float64 
 7   embarked      889 non-null    object  
 8   class         891 non-null    category
 9   who           891 non-null    object  
10  adult_male    891 non-null    bool    
11  deck          203 non-null    category
12  embark_town   889 non-null    object  
13  alive         891 non-null    object  
14  alone         891 non-null    bool    
dtypes: bool(2), category(2), float64(2), int64(4), object(5)
memory usage: 80.7+ KB

```

```
In [5]: data.isnull().sum().sort_values(ascending=False)
```

```

Out[5]: deck          688
age             177
embarked         2
embark_town      2
survived         0
pclass          0
sex             0
sibsp           0
parch           0
fare            0
class           0
who             0
adult_male      0
alive           0
alone           0
dtype: int64

```

```
In [6]: round(data.isnull().sum() / len(data) * 100, 2).sort_values(ascending=False)
```

```
Out[6]: deck          77.22  
age           19.87  
embarked       0.22  
embark_town    0.22  
survived       0.00  
pclass         0.00  
sex            0.00  
sibsp          0.00  
parch          0.00  
fare           0.00  
class          0.00  
who            0.00  
adult_male     0.00  
alive          0.00  
alone          0.00  
dtype: float64
```

```
In [7]: # Import Libraries  
import pandas as pd  
import numpy as np  
import seaborn as sns  
# Load titanic dataset  
data = sns.load_dataset('titanic')  
  
# calculate missing values  
print("-----")  
print(f"Missing values in each column:\n{data.isnull().sum().sort_values(ascending=False)}")  
print("-----")  
print(f"Percentage of missing values in each column:\n{round(data.isnull().sum() / len(data) * 100, 2).sort_values(a
```

```
-----  
Missing values in each column:  
deck          688  
age           177  
embarked       2  
embark_town    2  
survived       0  
pclass        0  
sex           0  
sibsp         0  
parch         0  
fare          0  
class         0  
who           0  
adult_male    0  
alive         0  
alone         0  
dtype: int64  
-----  
Percentage of missing values in each column:  
deck          77.22  
age           19.87  
embarked       0.22  
embark_town    0.22  
survived       0.00  
pclass        0.00  
sex           0.00  
sibsp         0.00  
parch         0.00  
fare          0.00  
class         0.00  
who           0.00  
adult_male    0.00  
alive         0.00  
alone         0.00  
dtype: float64
```

```
In [8]: data.head()
```

Out[8]:

	survived	pclass	sex	age	sibsp	parch	fare	embarked	class	who	adult_male	deck	embark_town	alive	alone
0	0	3	male	22.0	1	0	7.2500	S	Third	man	True	NaN	Southampton	no	False
1	1	1	female	38.0	1	0	71.2833	C	First	woman	False	C	Cherbourg	yes	False
2	1	3	female	26.0	0	0	7.9250	S	Third	woman	False	NaN	Southampton	yes	True
3	1	1	female	35.0	1	0	53.1000	S	First	woman	False	C	Southampton	yes	False
4	0	3	male	35.0	0	0	8.0500	S	Third	man	True	NaN	Southampton	no	True

In [9]: `round(data['age'].mean(), 2)`

Out[9]: 29.7

In [10]: `data['age'].median()`

Out[10]: 28.0

In [11]: `# Mean of age to fill age missing values
data['age'] = data['age'].fillna(data['age'].median())

drop deck column
data.drop('deck', axis=1, inplace=True)
data.head()`

Out[11]:

	survived	pclass	sex	age	sibsp	parch	fare	embarked	class	who	adult_male	embark_town	alive	alone
0	0	3	male	22.0	1	0	7.2500	S	Third	man	True	Southampton	no	False
1	1	1	female	38.0	1	0	71.2833	C	First	woman	False	Cherbourg	yes	False
2	1	3	female	26.0	0	0	7.9250	S	Third	woman	False	Southampton	yes	True
3	1	1	female	35.0	1	0	53.1000	S	First	woman	False	Southampton	yes	False
4	0	3	male	35.0	0	0	8.0500	S	Third	man	True	Southampton	no	True

```
In [12]: data['embark_town'].value_counts()
```

```
Out[12]: embark_town  
Southampton    644  
Cherbourg       168  
Queenstown      77  
Name: count, dtype: int64
```

```
In [13]: # replacing embarked missing values with mode  
data['embark_town'] = data['embark_town'].fillna(data['embark_town'].mode()[0])  
data['embarked'] = data['embarked'].fillna(data['embarked'].mode()[0])  
  
data.isnull().sum().sort_values(ascending=False)
```

```
Out[13]: survived      0  
pclass      0  
sex          0  
age          0  
sibsp        0  
parch        0  
fare         0  
embarked     0  
class        0  
who          0  
adult_male   0  
embark_town  0  
alive        0  
alone        0  
dtype: int64
```

we can also impute using sk learn

```
In [14]: # import libraries  
import pandas as pd  
import numpy as np  
import seaborn as sns  
  
from sklearn.impute import SimpleImputer
```

```
In [15]: df = sns.load_dataset('titanic')

# impute age column using simpleimputer from sklearn
imputer = SimpleImputer(strategy='mean')

df['age'] = imputer.fit_transform(df[['age']])
df.isnull().sum().sort_values(ascending=False)
```

```
Out[15]: deck          688
embarked          2
embark_town       2
survived          0
pclass           0
sex              0
age              0
sibsp            0
parch            0
fare             0
class            0
who              0
adult_male       0
alive            0
alone            0
dtype: int64
```

```
In [16]: df.head()
```

```
Out[16]:
```

	survived	pclass	sex	age	sibsp	parch	fare	embarked	class	who	adult_male	deck	embark_town	alive	alone
0	0	3	male	22.0	1	0	7.2500	S	Third	man	True	NaN	Southampton	no	False
1	1	1	female	38.0	1	0	71.2833	C	First	woman	False	C	Cherbourg	yes	False
2	1	3	female	26.0	0	0	7.9250	S	Third	woman	False	NaN	Southampton	yes	True
3	1	1	female	35.0	1	0	53.1000	S	First	woman	False	C	Southampton	yes	False
4	0	3	male	35.0	0	0	8.0500	S	Third	man	True	NaN	Southampton	no	True

multivariate imputation

```
In [17]: df = sns.load_dataset('titanic')

from sklearn.experimental import enable_iterative_imputer
from sklearn.impute import IterativeImputer

# impute age column using iterativeimputer from sklearn
imputer = IterativeImputer(max_iter=20, n_nearest_features=5)
df['age'] = imputer.fit_transform(df[['age']])

df.isnull().sum().sort_values(ascending=False)
```

```
Out[17]: deck          688
embarked          2
embark_town        2
survived           0
pclass            0
sex               0
age              0
sibsp            0
parch            0
fare             0
class            0
who              0
adult_male        0
alive            0
alone            0
dtype: int64
```

ffill and backward fill

```
In [18]: df = sns.load_dataset('titanic')
```

```
df.isnull().sum().sort_values(ascending=False)
```

```
Out[18]: deck          688  
age            177  
embarked       2  
embark_town    2  
survived       0  
pclass        0  
sex            0  
sibsp         0  
parch         0  
fare          0  
class         0  
who           0  
adult_male    0  
alive         0  
alone         0  
dtype: int64
```

```
In [19]: # using forward fill impute age column  
df['age'] = df['age'].bfill()
```

using KNN imputer

```
In [20]: from sklearn.impute import KNNImputer  
  
# impute age column using KNNImputer from sklearn  
imputer = KNNImputer(n_neighbors=5)  
df['age'] = imputer.fit_transform(df[['age']])  
  
# drop rows having missing values  
df.dropna(inplace=True)  
  
df.info()
```

```

<class 'pandas.core.frame.DataFrame'>
Index: 201 entries, 1 to 889
Data columns (total 15 columns):
 #   Column      Non-Null Count  Dtype
---  -
 0   survived    201 non-null    int64
 1   pclass      201 non-null    int64
 2   sex         201 non-null    object
 3   age         201 non-null    float64
 4   sibsp       201 non-null    int64
 5   parch       201 non-null    int64
 6   fare        201 non-null    float64
 7   embarked    201 non-null    object
 8   class       201 non-null    category
 9   who         201 non-null    object
10  adult_male   201 non-null    bool
11  deck        201 non-null    category
12  embark_town  201 non-null    object
13  alive       201 non-null    object
14  alone       201 non-null    bool
dtypes: bool(2), category(2), float64(2), int64(4), object(5)
memory usage: 20.1+ KB

```

```
In [21]: df.isnull().sum().sort_values(ascending=False)
```

```

Out[21]: survived    0
pclass      0
sex          0
age          0
sibsp        0
parch        0
fare         0
embarked     0
class        0
who          0
adult_male   0
deck         0
embark_town  0
alive        0
alone        0
dtype: int64

```

In []: